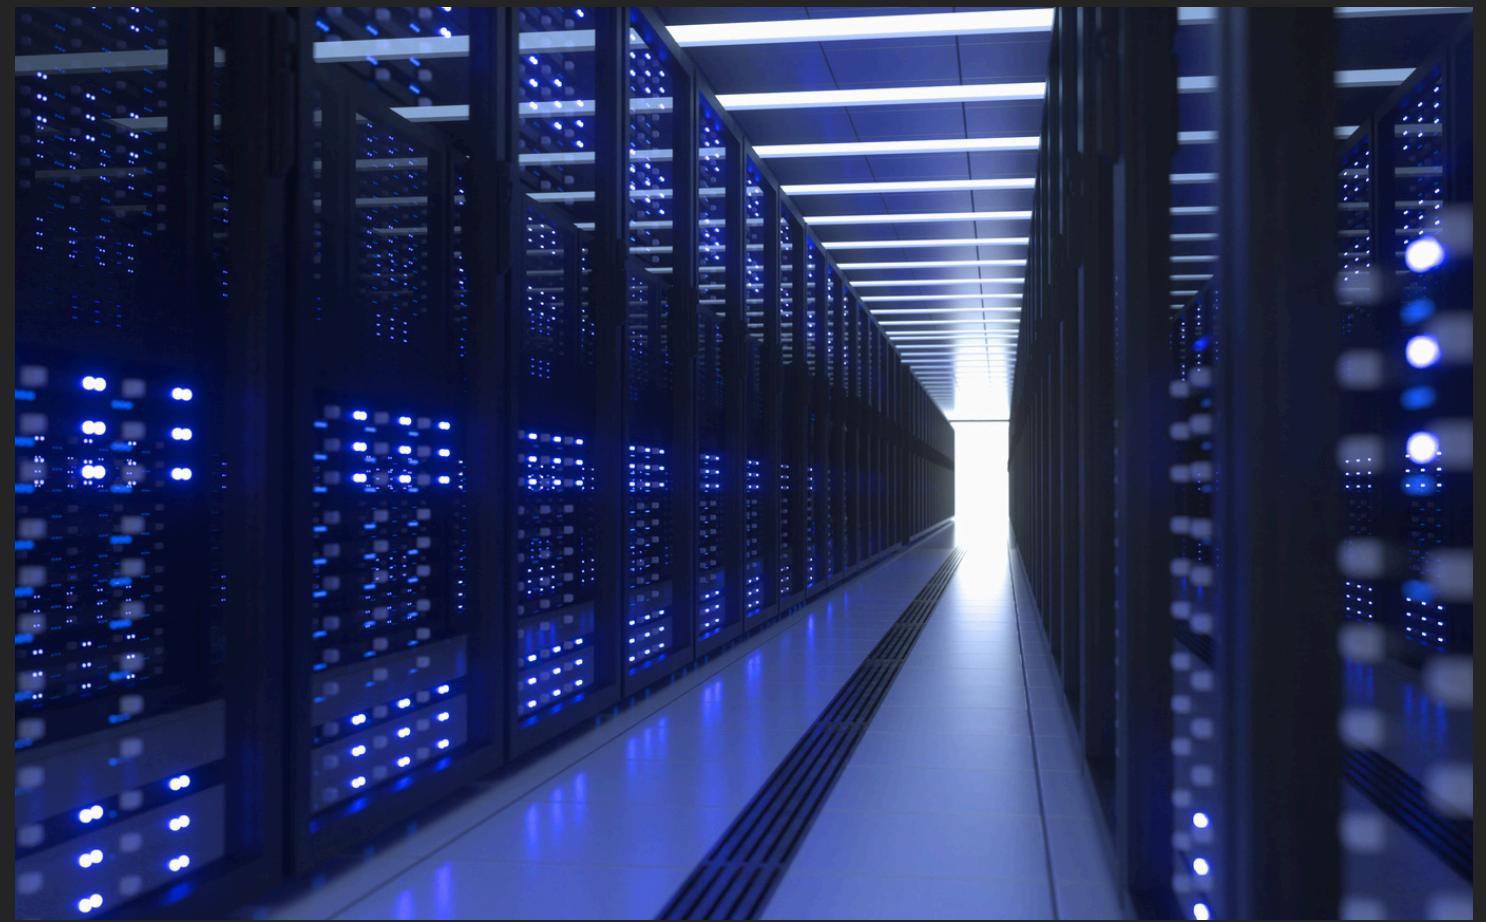




# 2024 APAC

# HPC-AI

Universiti Putra Malaysia  
UPMTeam3



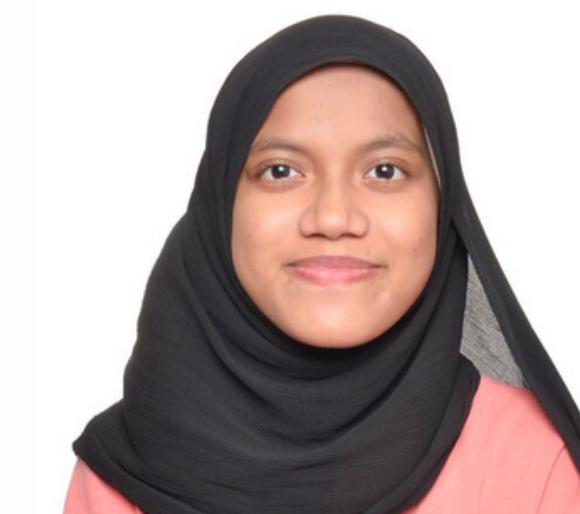
# Our Team



**Anis Azman**



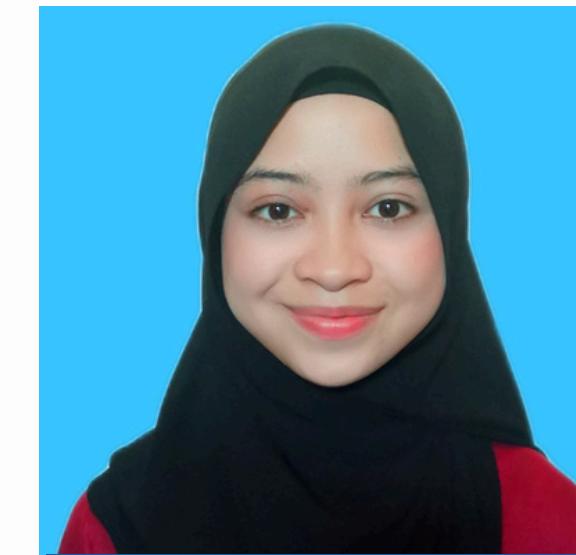
**Nurul Farizatul Aina**



**Maimunah Hosni**



**Wan Siti Aisyah**



**Azyan Syazwani**



**Siti Nurinsyirah**



# Universiti Putra Malaysia

Universiti Putra Malaysia (UPM) is a leading Malaysian university renowned for its expertise in agriculture and related fields. Founded in 1931, UPM has a long history of academic excellence and has grown to become a global leader in its area of expertise.

UPM offers a diverse range of programs, including food crops, animal husbandry, veterinary medicine, and forestry. The university also has expanded its scope to include science, technology, and other disciplines relevant to society and the national needs, ensuring graduates are well-prepared for the challenges of the modern world. UPM is committed to environmental sustainability and has consistently ranked highly in green university rankings.

# OUR SUPERVISORS



**MRS. Y.M. RAJA AZLINA  
BINTI RAJA MAHMOOD**



**PROF. MADYA DR. NOR ASILAH  
WATI BINTI ABDUL HAMID**

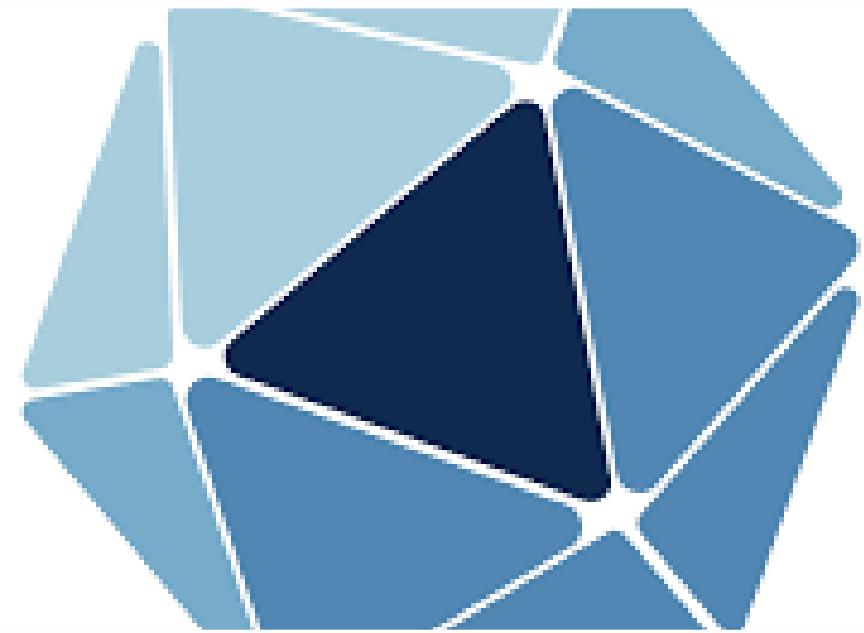


# Faculty of Computer Science and Information Technology (FCSIT)

In the 1980s, UPM started offering computer science courses to meet growing demand. By 1992, a dedicated department was established, offering various courses, including postgraduate programs.

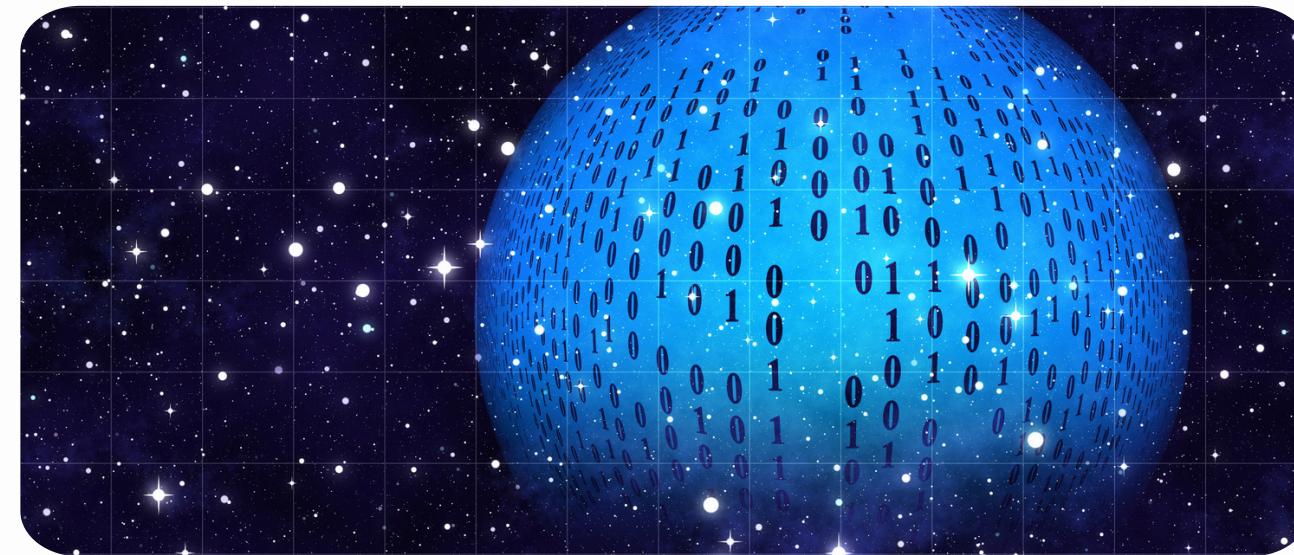
In 1998, the Faculty of Computer Science and Information Technology was created to meet the increasing demand for computer education and keep up with technological advancements. Today, the faculty has four departments: Multimedia, Computer Science, Software Engineering & Information System, and Communication Technology & Networks.

# Definitions



## HOOMD-BLUE

HOOMD-blue is a molecular dynamics simulation tool used to study soft matter systems like polymers and colloids. It's designed to be easy to use, efficient, and flexible, making it a valuable resource for researchers in many fields.



## LITGPT LLaMA2

Lit-GPT is an open-source project that provides a lightweight and optimized version of large language models, like LLaMA 2, developed by Meta AI. LLaMA 2 is a powerful model used for tasks such as text generation and summarization.

Lit-GPT can run on regular computers, allowing researchers and developers to use them without needing powerful hardware.

# Our Methodology

Improve performance, ideas, and tuning methods.



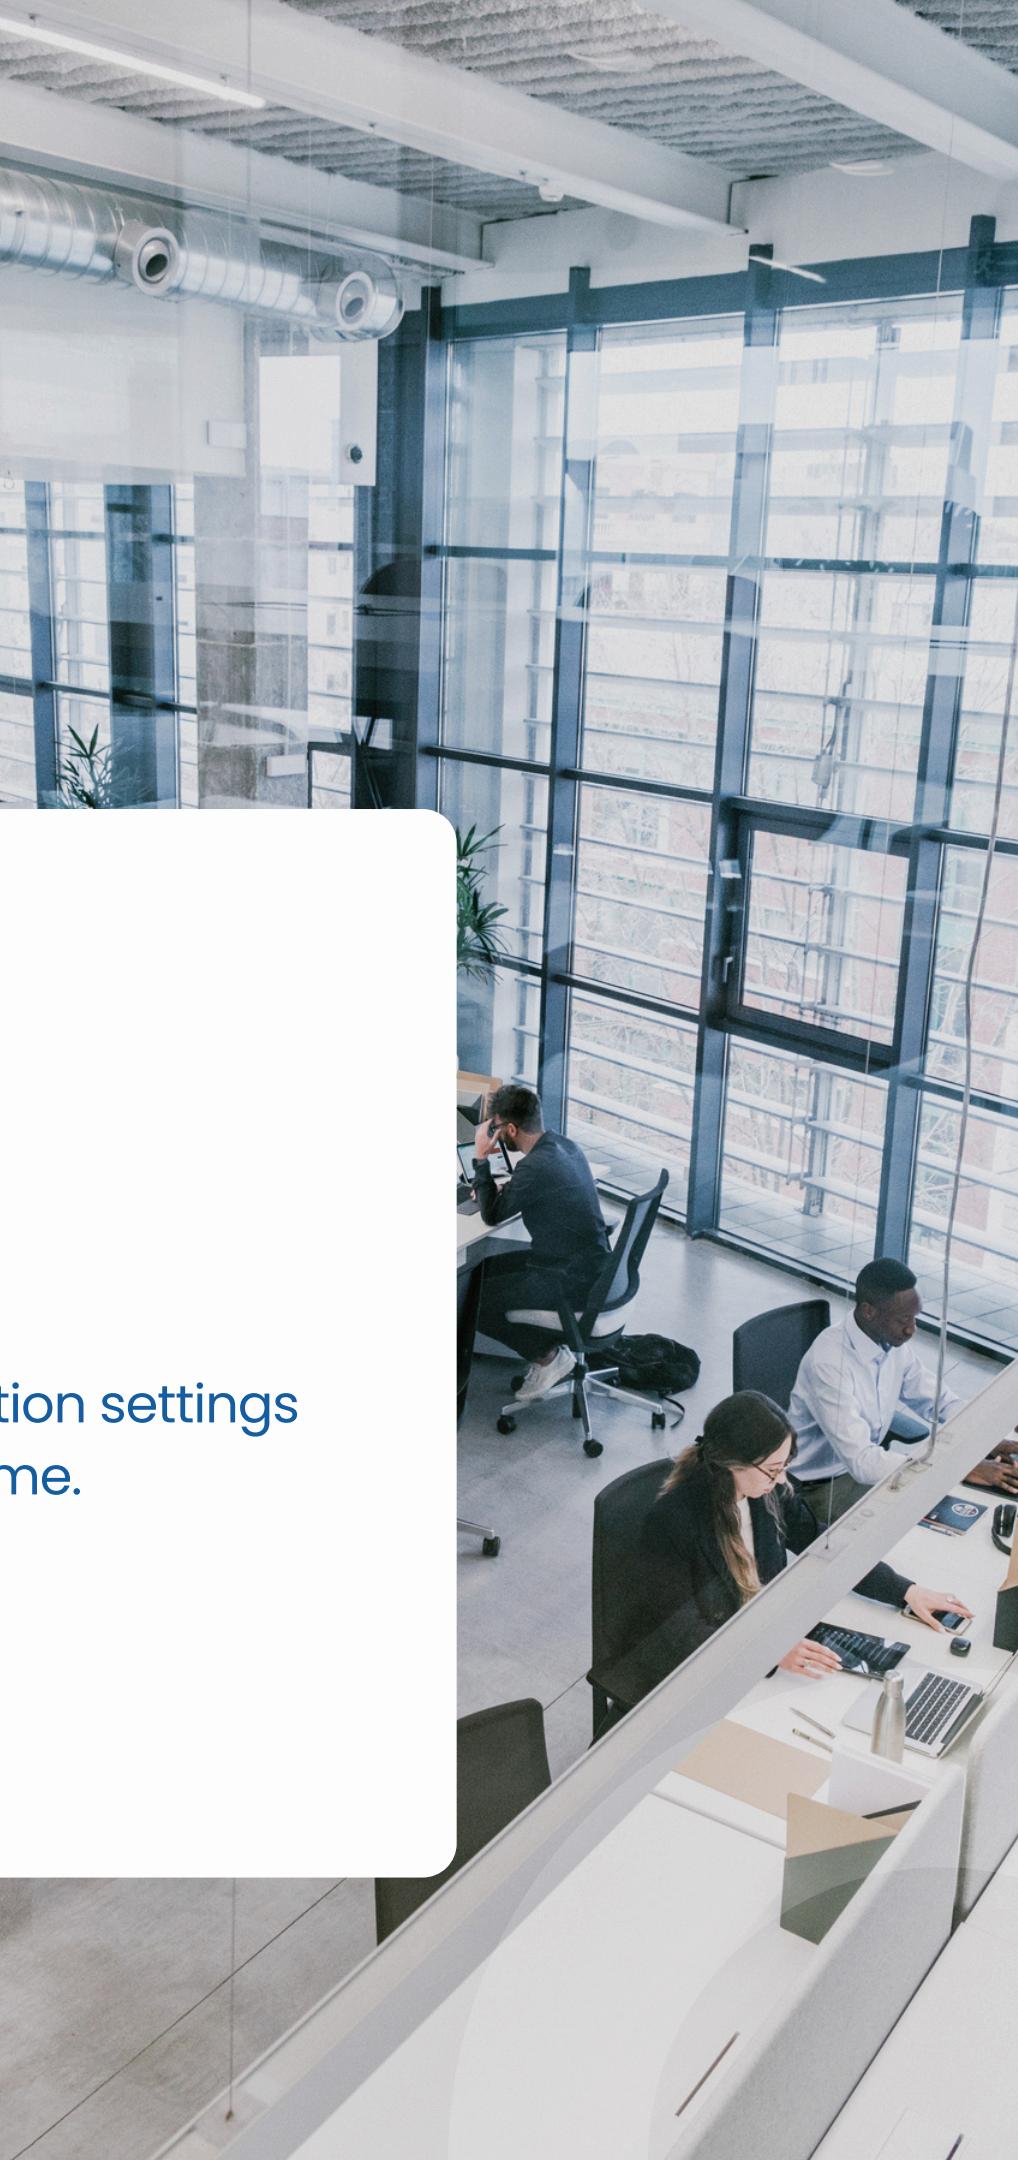
## HOOMDBLUE

- Calculate efficiency using total cores and speedup.
- Try allocating different cores and nodes to increase the number of steps.
- Result: Comparison on different number of nodes.



## LLAMA2

- Optimization on MPI communication settings
- Result: Comparison on training time.
- Finetune using LoRA



# **TASK 1**

## **HOOMD-BLUE USING GADI**

# Performance Metrics

- 01 **Steps per second** – This measures how fast the system can simulate particle movements over time.
- 02 **Execution time** – This shows how long the task takes to complete.
- 03 **Speedup** – This measures how much faster a task completes when multiple processors are used compared to using a single processor.
- 04 **Efficiency** – This measures how effectively the processors are being used in parallel.

# PBS Script Directives & Modules Loaded

```
#!/bin/bash
#PBS -j oe
#PBS -M 215014@student.upm.edu.my,214511@student.upm.edu.my
#PBS -m abe
#PBS -P zd64
#PBS -l ngpus=0
#PBS -l walltime=00:00::60
#PBS -l other=hyperthread
#-report-bindings \
module purge
module load ${HOME}/hpcx-v2.20-gcc-mlnx_ofed-redhat8-cuda12-x86_64/modulefiles/hpcx-ompi
```

hoomd.sh

# PBS Commands Used

- qstat - To display queue information
- qdel - To delete a job
- cat <job file> - To display the contents of a file

```
(/home/552/sp1115/scratch/workdir/hoomd/hoomd.py312) [sp1115@gadi-login-05 run]$ qstat
Job id                               Name           User           Time Use S Queue
-----                               -----           -----           ----- - -----
126480181.gadi-pbs      hoomd.nodes1.WS* sp1115          00:10:27 R normal-exec
```

```
[sp1115@gadi-login-05 run]$ qdel 126480181.gadi-pbs
```

```
(base) [sp1115@gadi-login-07 run]$ cat hoomd.nodes4.WS40000.BS80000.o126485280
time mpirun -host gadi-cpu-clx-2378.gadi.nci.org.au,gadi-cpu-clx-2379.gadi.nci.org.au,gadi-cpu-clx-2380.gadi.nci.or
adi-cpu-clx-2391.gadi.nci.org.au,gadi-cpu-clx-2396.gadi.nci.org.au,gadi-cpu-clx-2400.gadi.nci.org.au,gadi-cpu-clx-2
di.nci.org.au,gadi-cpu-clx-2411.gadi.nci.org.au -wdir /home/552/sp1115/scratch/workdir/hoomd -output-filename /home
p-by ppr:48:node -oversubscribe -use-hwthread-cpus -x PYTHONPATH=/home/552/sp1115/scratch/workdir/hoomd/build/hoomd
e/552/sp1115/scratch/workdir/hoomd/hoomd.py312/bin/python -m hoomd_benchmarks.md_pair_wca --device CPU -v -N 200000
Using existing initial_configuration_cache/hard_sphere_200000_1.0_3.gsd
notice(2): Using domain decomposition: n_x = 8 n_y = 8 n_z = 9.
Running MDPairWCA benchmark
.. warming up for 40000 steps
.. running for 80000 steps 1 time(s)
.. 4240.476353911216 time steps per second
4240.476353911216
1243.40user 278.41system 0:37.44elapsed 4063%CPU (0avgtext+0avgdata 441320maxresident)k
347488inputs+16840outputs (39780major+8426785minor)pagefaults 0swaps

=====
Resource Usage on 2024-10-09 19:06:01:
Job Id: 126485280.gadi-pbs
Project: zd64
Exit Status: 0
Service Units: 12.80
NCPUs Requested: 576
NCPUs Used: 576
CPU Time Used: 05:03:03
Memory Requested: 576.0GB
Memory Used: 187.65GB
Walltime requested: 00:10:00
Walltime Used: 00:00:40
JobFS requested: 1.17GB
JobFS used: 0B
=====
```

# PBS Value Initialization

<b>Number of nodes</b>	<b>Number of cores used</b>	<b>Warmup/Benchmark</b>	<b>Walltime Requested</b>	<b>Memory Requested</b>
1 x 2	48 x 1	40,000/80,000	10 mins	48GB
2 x 2	48 x 2	40,000/80,000	10 mins	96GB
4 x 2	48 x 4	40,000/80,000	10 mins	192GB
8 x 2	48 x 8	40,000/80,000	10 mins	384GB
16 x 2	48 x 16	10,000/160,000	10 mins	768GB
32 x 2	48 x 32	10,000/320,000	10 mins	1536GB

# Execution Script

```
module purge
module load ${HOME}/hpcx-v2.20-gcc-mlnx_ofed-redhat8-cuda12-x86_64/modulefiles/hpcx-ompi

hosts=$(sort -u ${PBS_NODEFILE} | paste -sd ',')

cmd="time mpirun \
-host ${hosts} \
-wdir ${HOME}/scratch/workdir/hoomd \
-output-filename ${HOME}/run/output/${PBS_JOBNAME}.${PBS_JOBID} \
-map-by ppr:$((1*${NCPUS})):node \
-oversubscribe -use-hwthread-cpus \
-x PYTHONPATH=${HOME}/scratch/workdir/hoomd/build/hoomd-openmpi4.1.5:${HOME}/scratch/workdir/hoomd/hoomd-
benchmarks \
${HOME}/scratch/workdir/hoomd/hoomd.py312/bin/python \
-m hoomd_benchmarks.md_pair_wca \
--device CPU -v \
-N ${N} --repeat ${repeat} \
--warmup_steps ${warmup_steps} --benchmark_steps ${benchmark_steps}"

echo ${cmd}

exec ${cmd}
~
```

hoomd.sh

# Submitting Job Commands

## Example:

```
(base) [sp1115@gadi-login-03 run]$ nodes=32 walltime=00:10:00 \
> warmup_steps=10000 benchmark_steps=8000 repeat=1 N=200000 \
\
'qsub -V \
-l> bash -c \
> 'qsub -V \
> -l walltime=${walltime},ncpus=$((48*nodes)),mem=$((48*nodes*1))gb \
> -N hoomd.nodes${nodes}.WS${warmup_steps}.BS${benchmark_steps} \
> hoomd.sh'
```

# Examples of Output

```
Using existing initial_configuration_cache/hard_sphere_200000_1.0_3.gsd
notice(2): Using domain decomposition: n_x = 4 n_y = 4 n_z = 6.
Running MDPairWCA benchmark
.. warming up for 40000 steps
.. running for 80000 steps 1 time(s)
.. 999.3391245451993 time steps per second
999.3391245451993
5598.51user 308.30system 2:08.09elapsed 4611%CPU (0avgtext+0avgdata 431800maxresident)k
225936inputs+16848outputs (28009major+12530102minor)pagefaults 0swaps

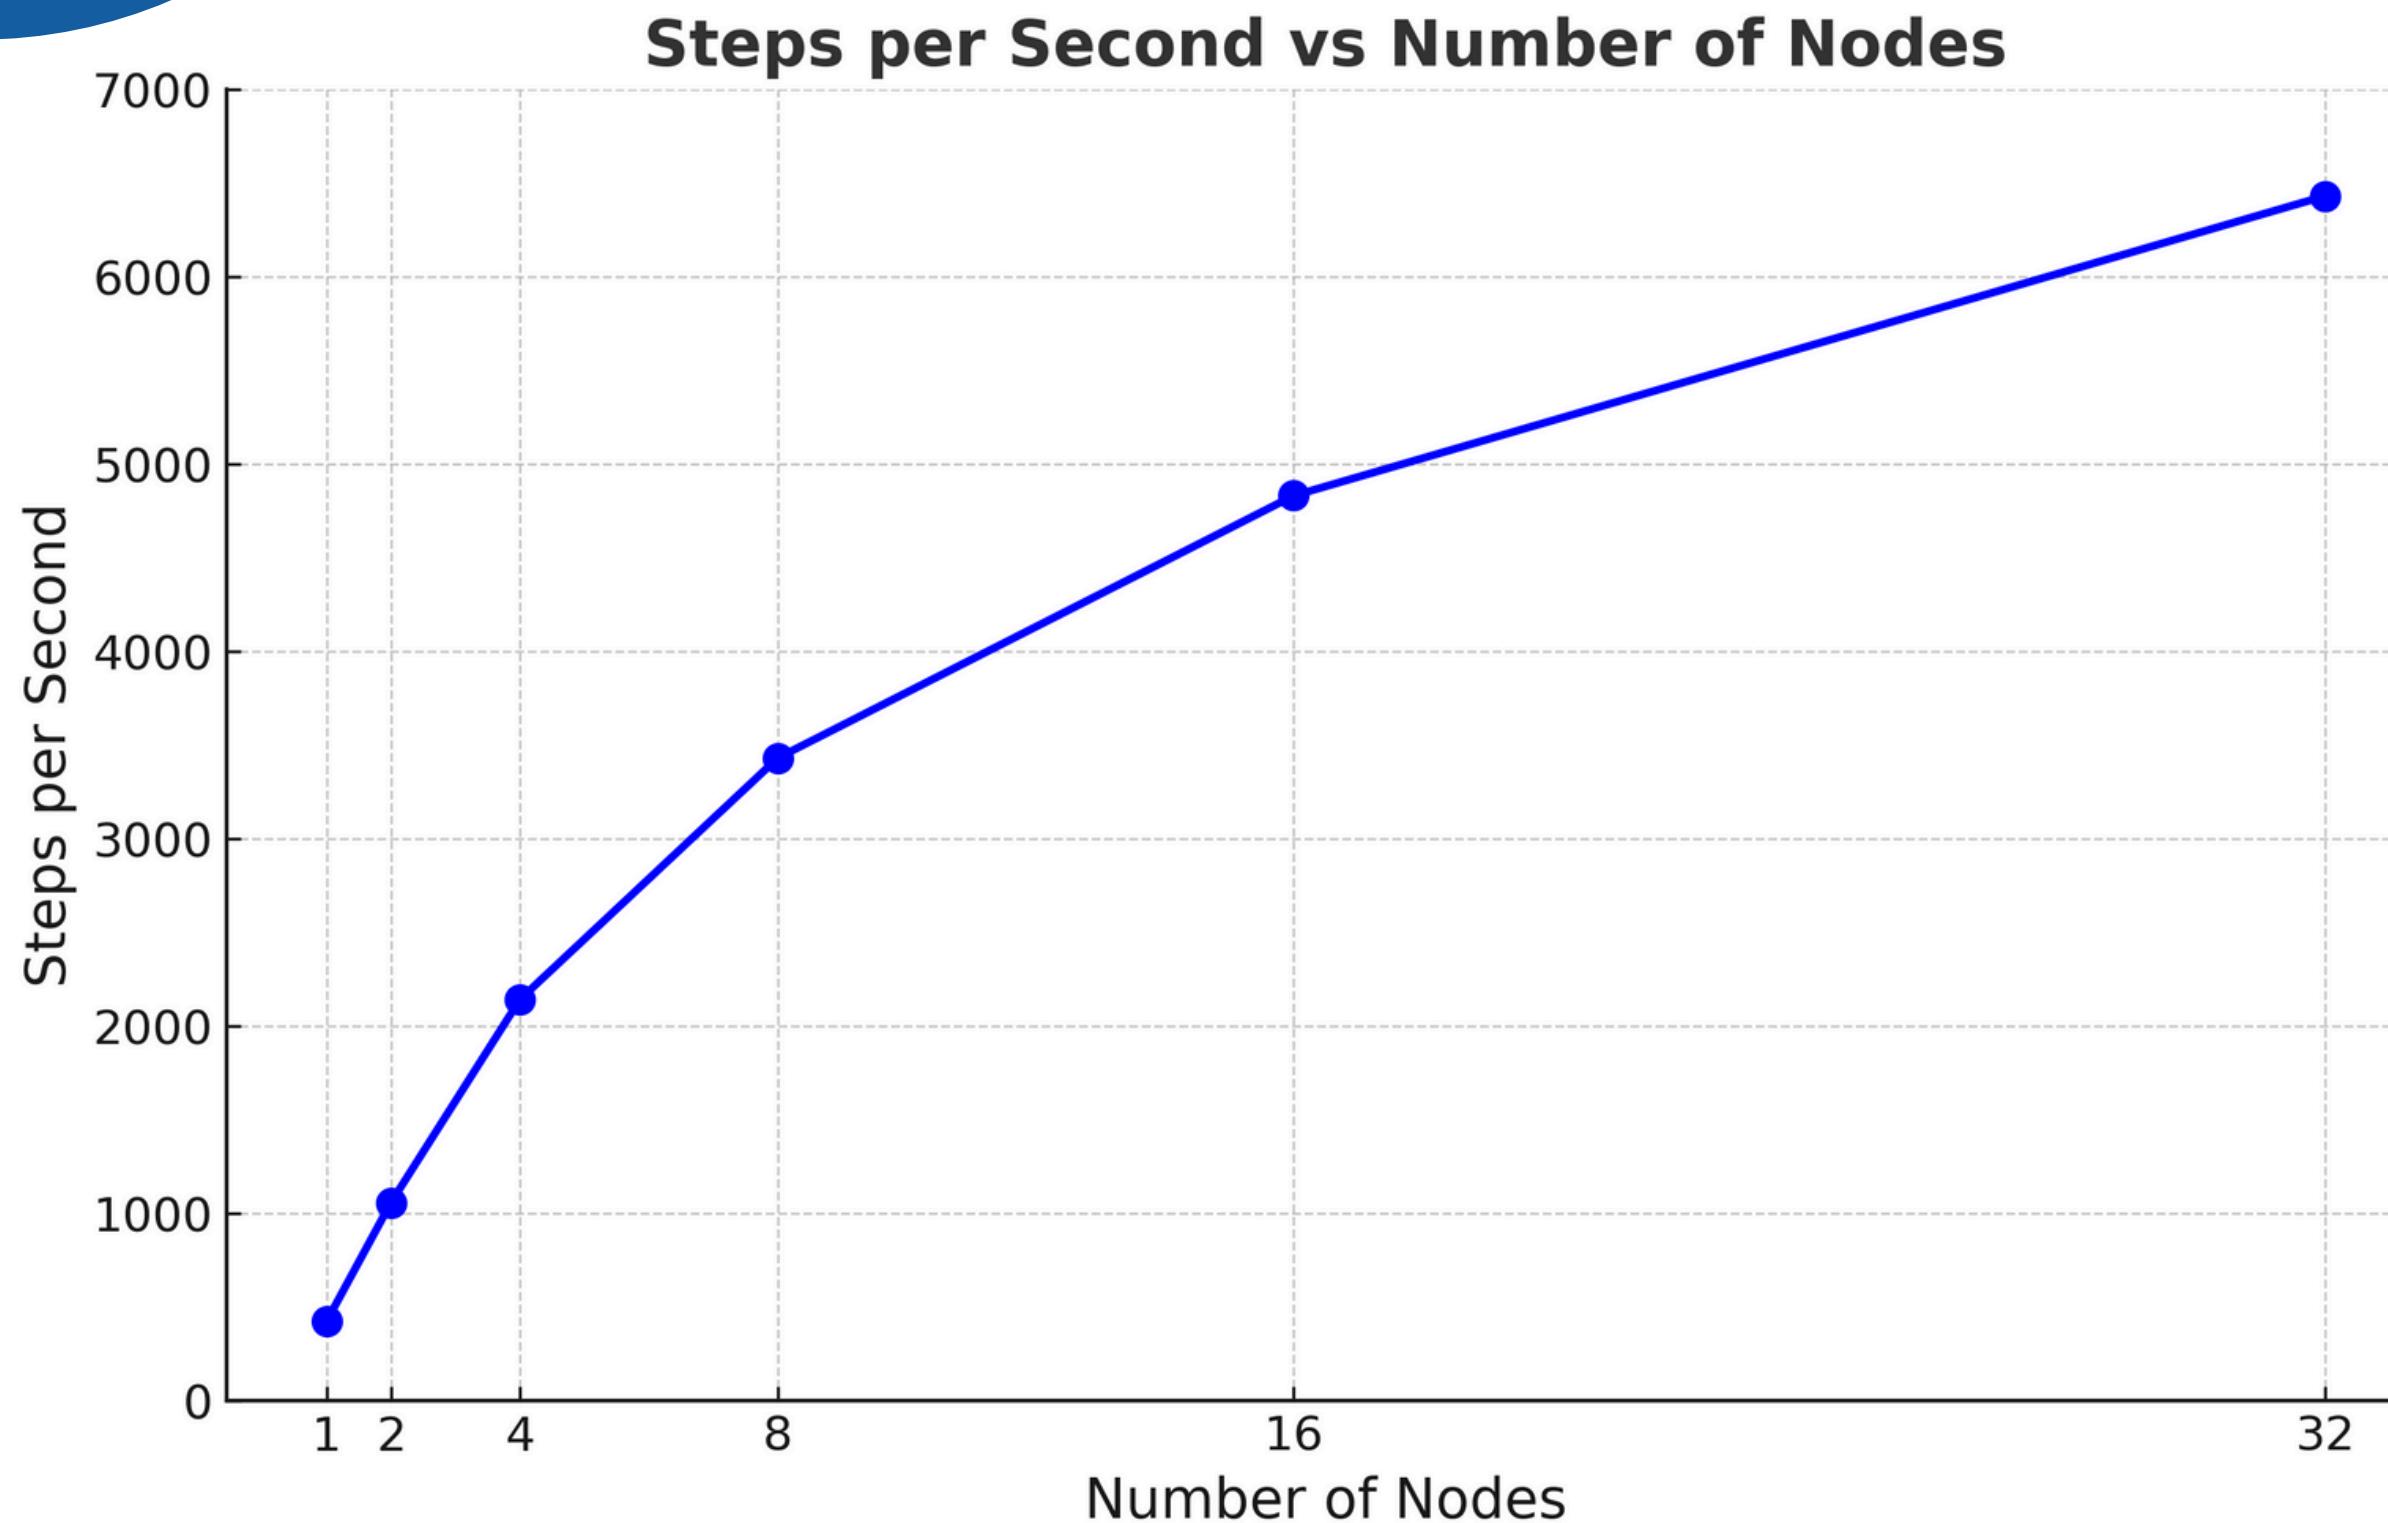
=====
Resource Usage on 2024-10-10 01:31:28:
Job Id: 126506621.gadi-pbs
Project: zd64
Exit Status: 0
Service Units: 6.93
NCPUUs Requested: 96 NCPUs Used: 96
CPU Time Used: 03:16:35
Memory Requested: 48.0GB Memory Used: 31.53GB
Walltime requested: 00:10:00 Walltime Used: 00:02:10
JobFS requested: 200.0MB JobFS used: 0B
=====
```

hoomd.sh output file

# HOOAMDBLUE Result

<b>Number of nodes</b>	<b>Number of cores used</b>	<b>Total Cores</b>	<b>Memory requested (GB)</b>	<b>Walltime Used</b>	<b>Memory Used (GB)</b>	<b>Steps per second</b>
1 x 2	48 x 1	48	48	0:55	21.54	423
2 x 2	48 x 2	96	96	0:27	31.68	1058
4 x 2	48 x 4	192	192	0:18	62.11	2142
8 x 2	48 x 8	384	384	0:15	123.58	3431
16 x 2	48 x 16	768	768	0:14	249.8	4831
32 x 2	48 x 32	1536	1536	0:14	494.79	6431

# Steps per second

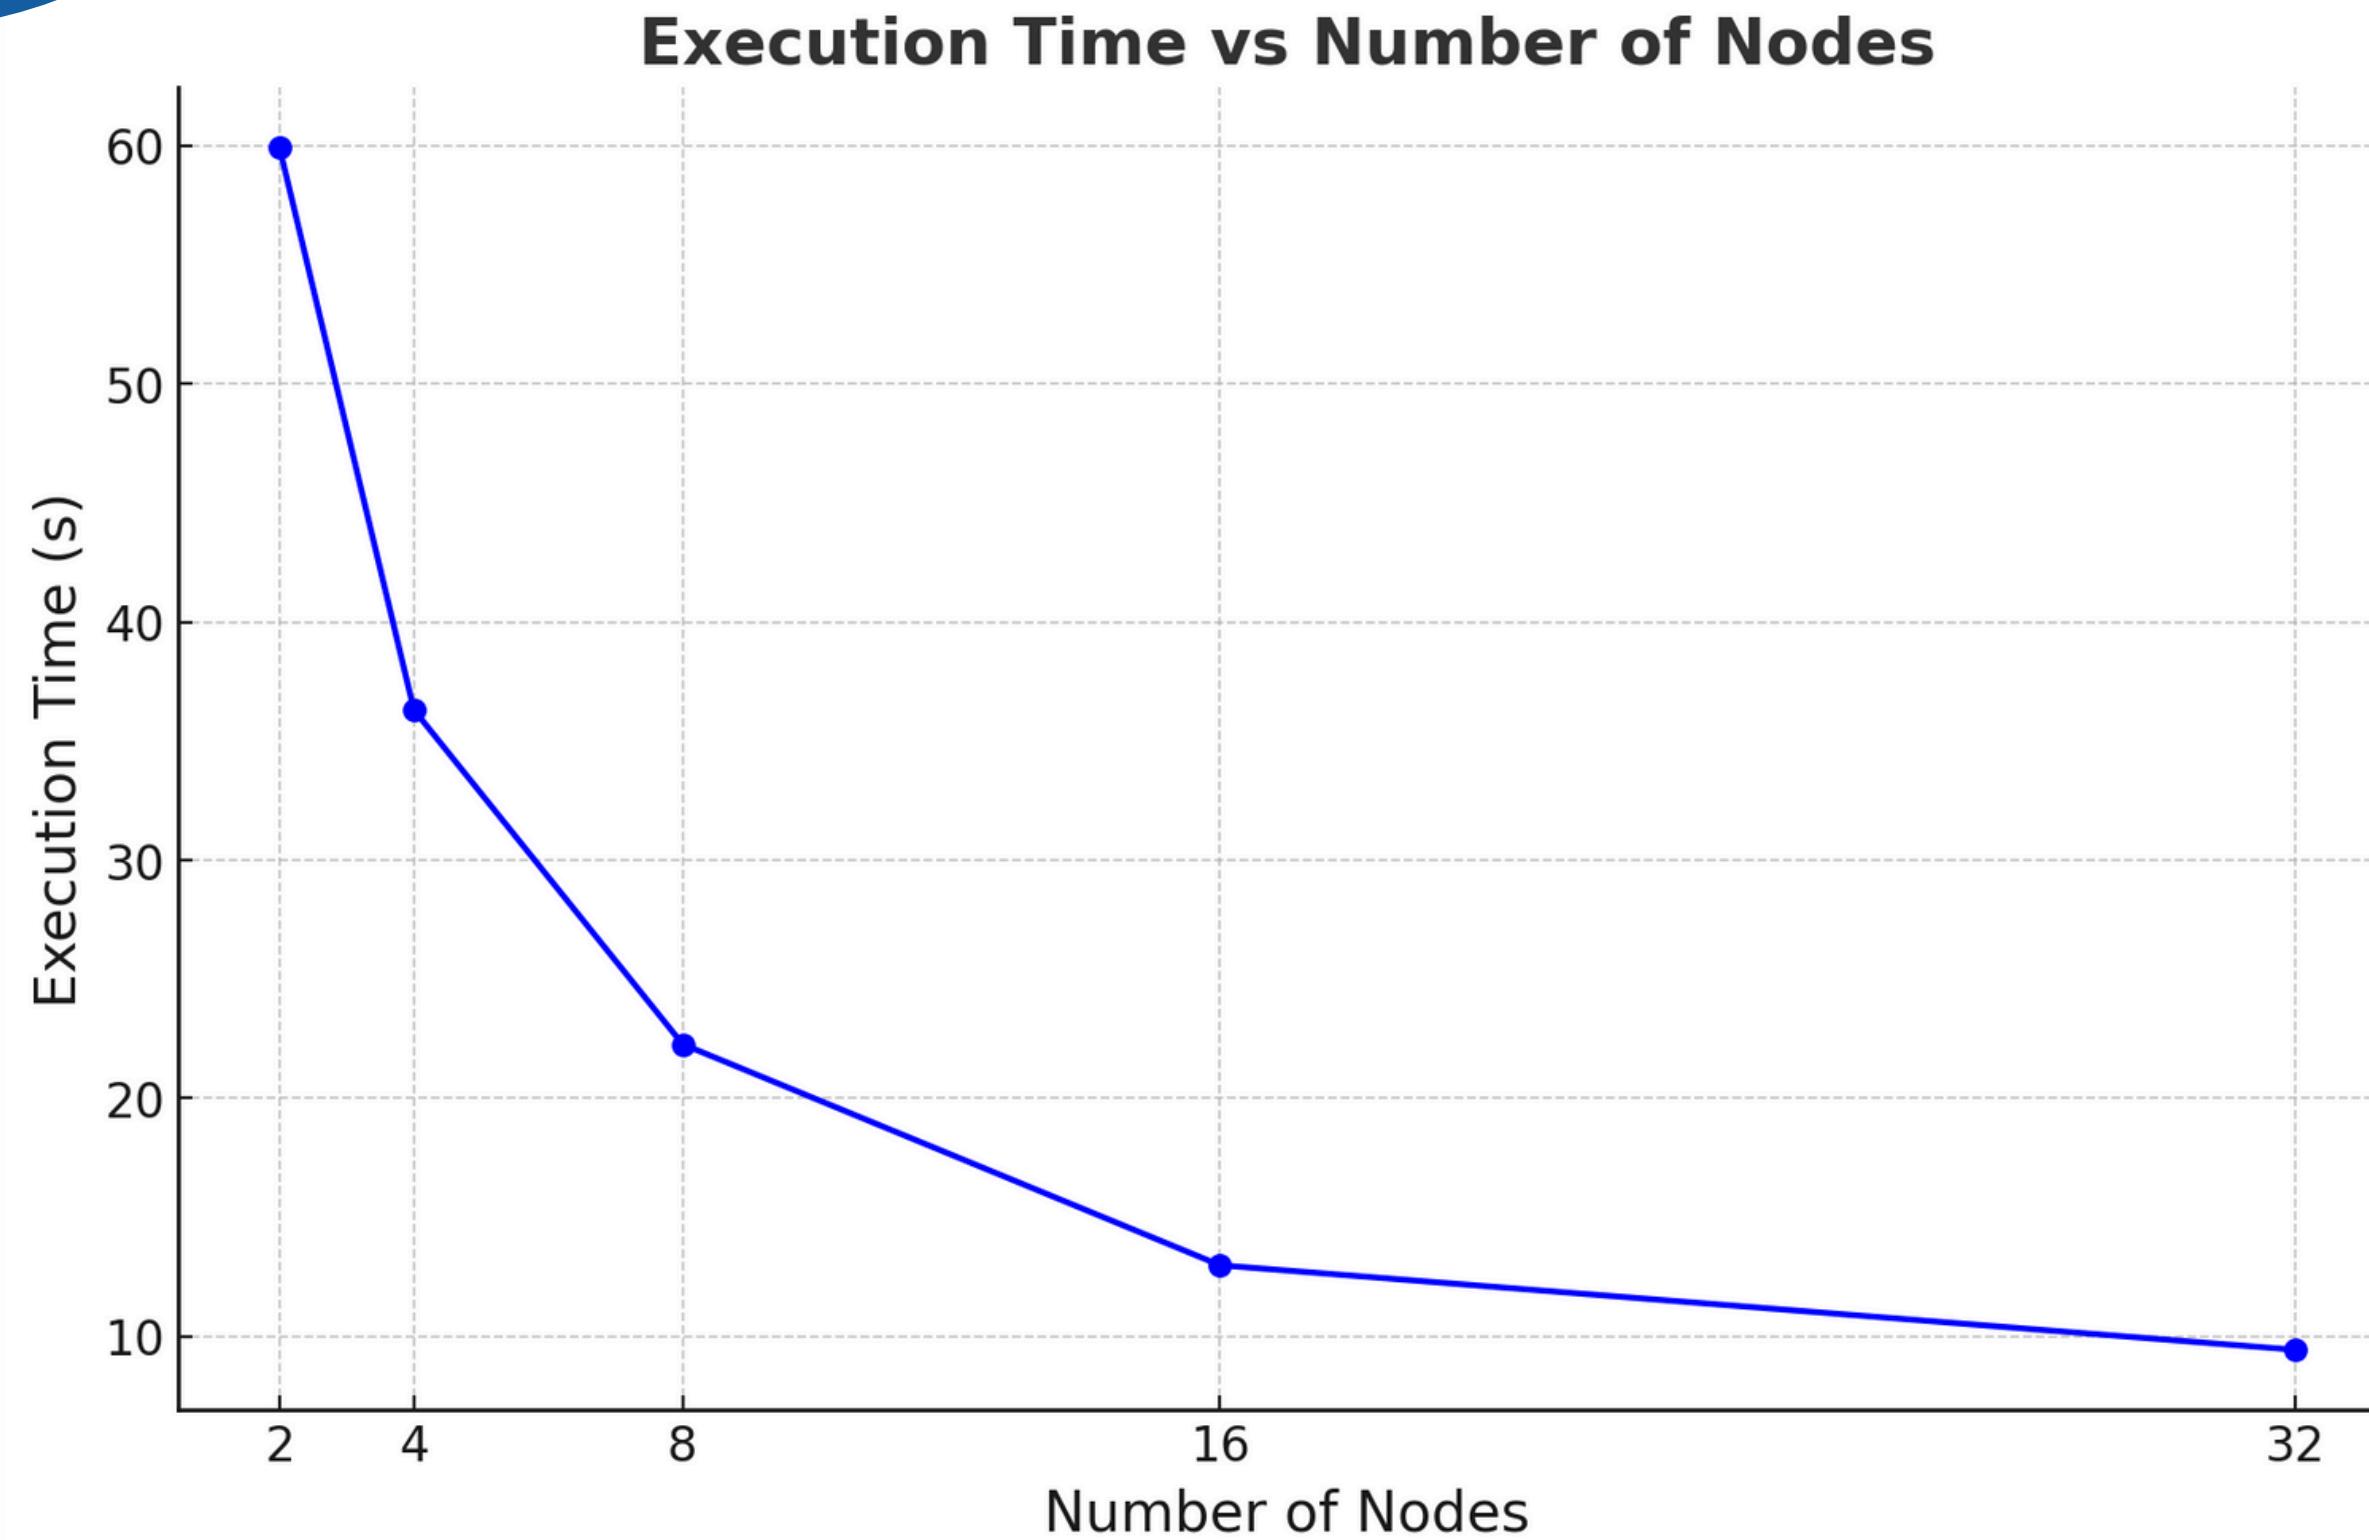


The steps per second increase as more nodes and cores are used, reflecting improved performance with additional computational resources.

However, while the performance improves significantly as nodes increase, the performance gain from adding more nodes diminishes slightly after 16 nodes.

Beyond this point, further node increases result in only marginal improvements, likely due to communication overheads and resource bottlenecks. This pattern suggests that adding nodes improves performance but with limited efficiency gains at higher node counts, which can guide optimal resource allocation.

# Execution Time(s)



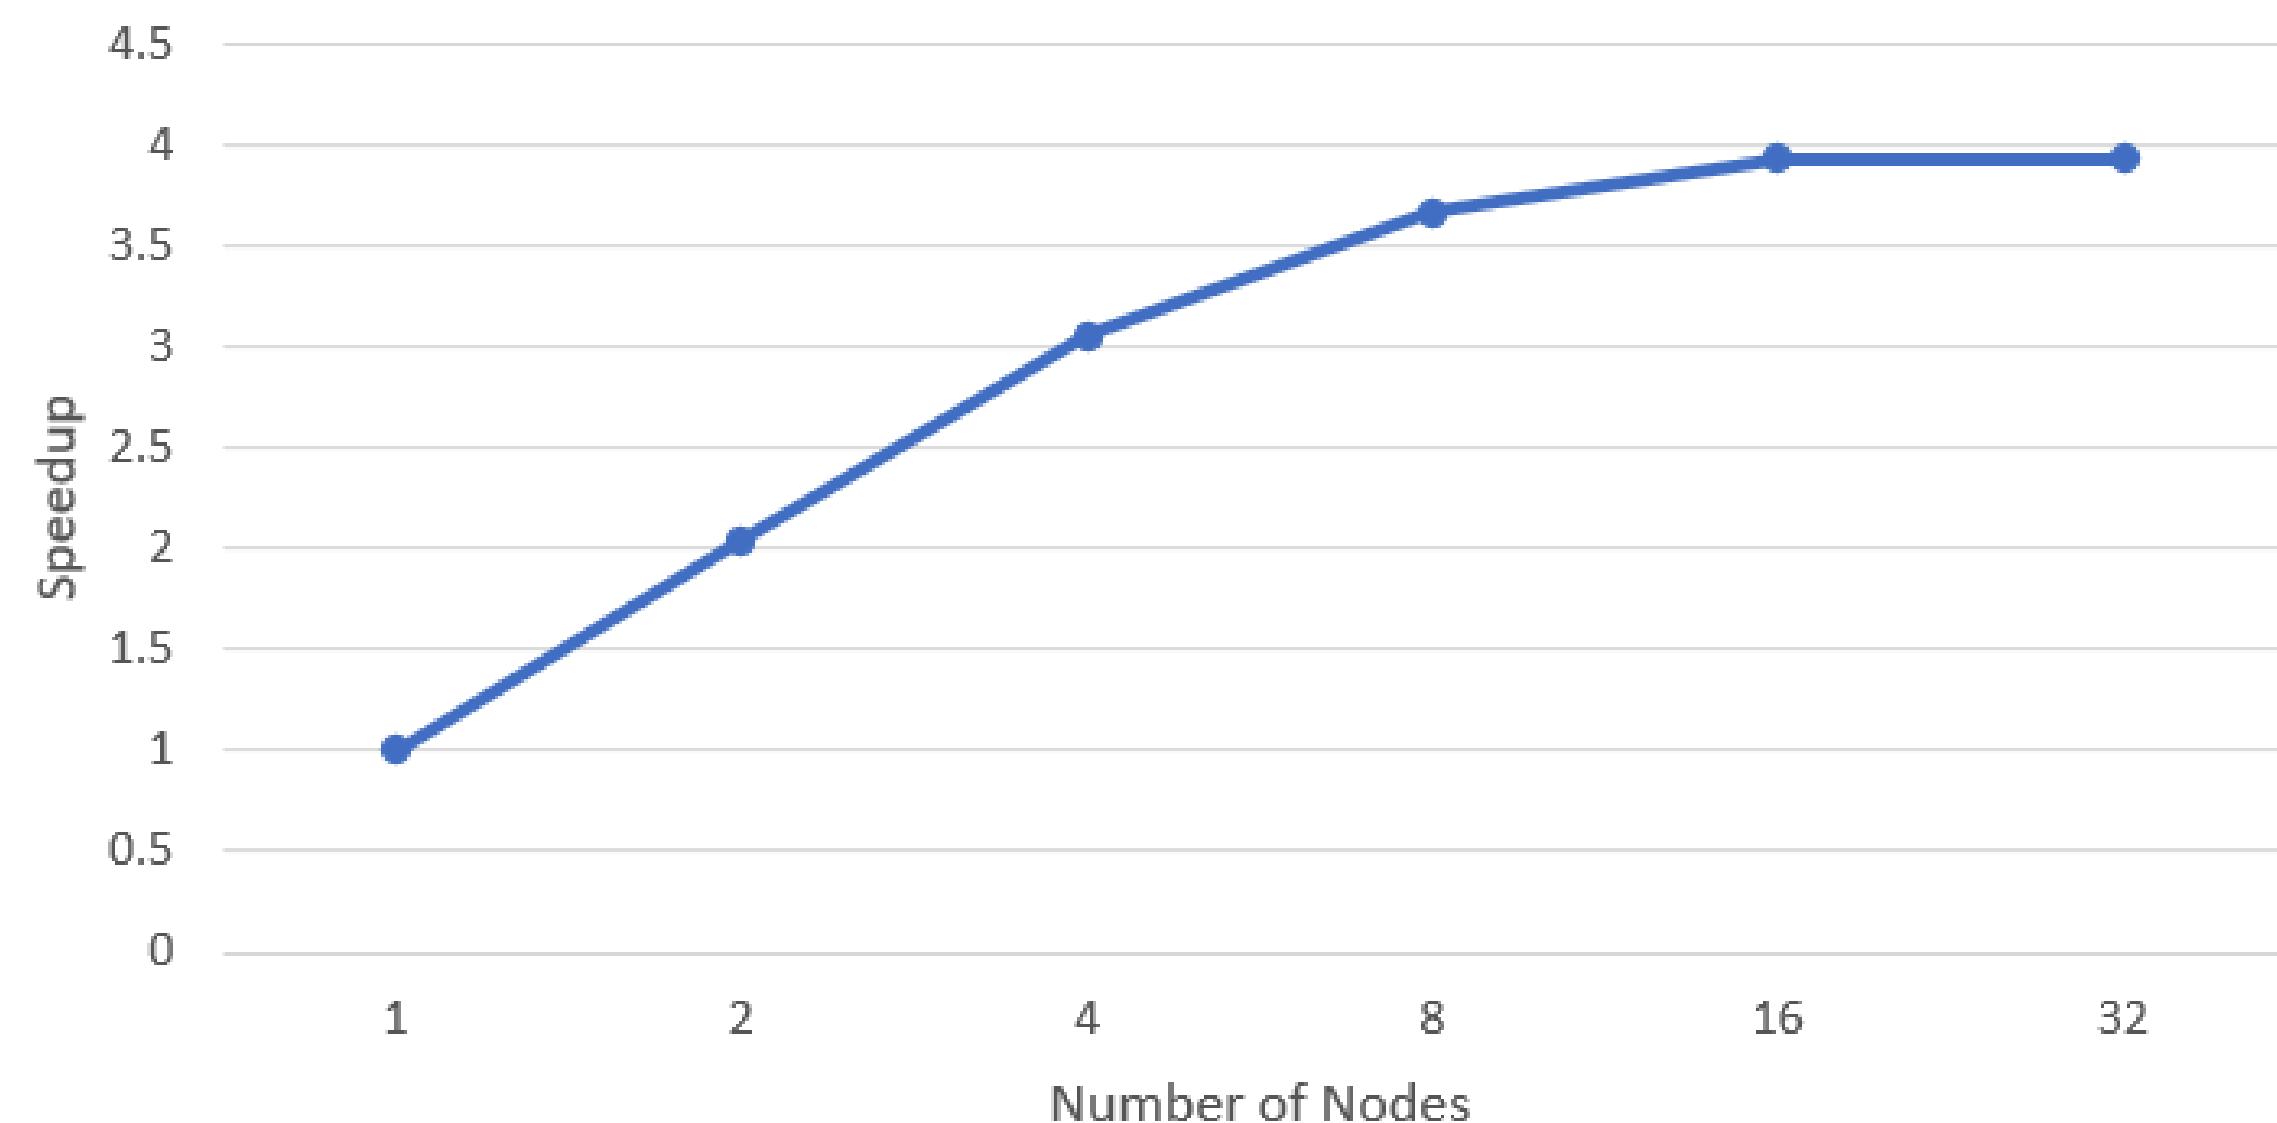
This graph shows that as the number of nodes increases, the walltime (execution time) decreases.

Adding more nodes allows the system to distribute the task across multiple nodes and process them simultaneously, which reduces the overall time required to complete the task.

However, beyond 16 nodes, we notice a slight increase in execution time, which could be due to the communication overhead between nodes, leading to delays.

# Simulation speed

Speedup vs Number of Nodes



As the total cores increases, the speedup of the system tends to increase because tasks are divided and processed in parallel, reducing overall time. This indicates that adding more cores can improve performance.

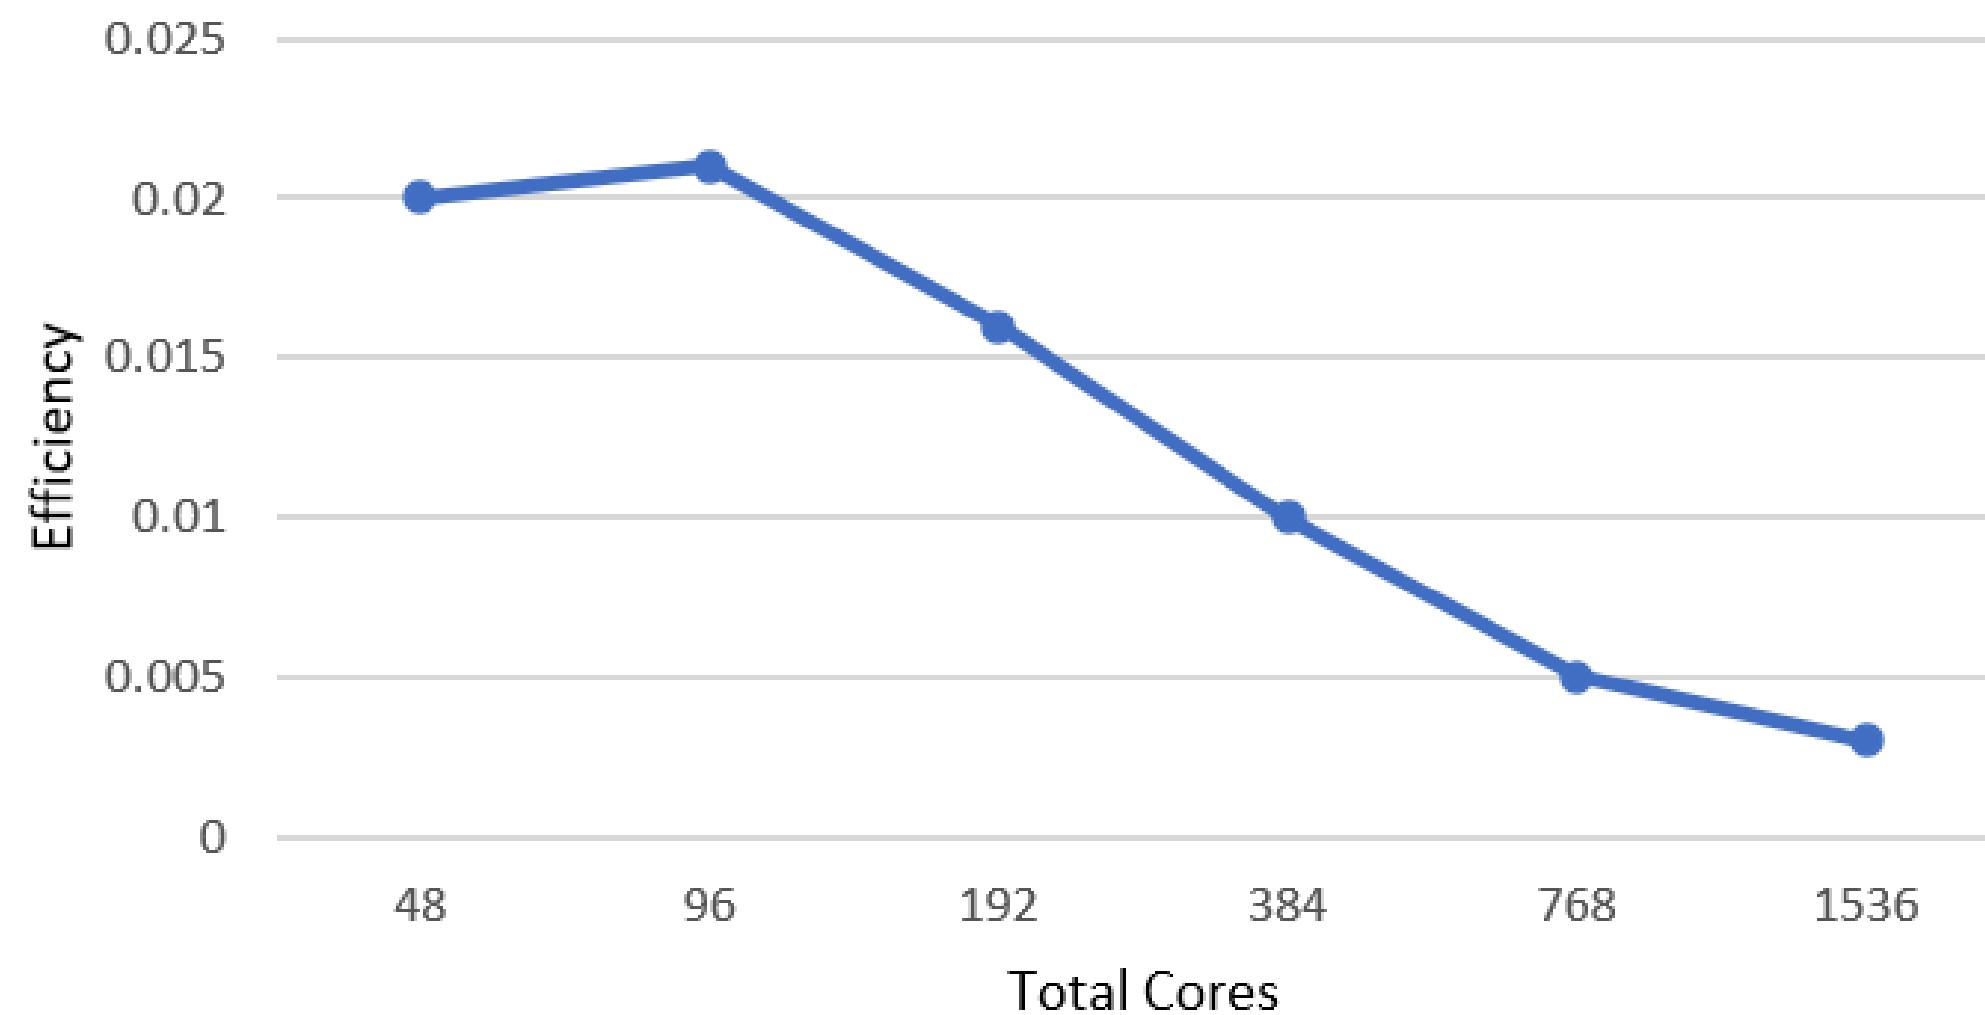
However, after a certain point, adding more cores offers diminishing returns due to communication overhead.

Number of nodes	1	2	4	8	16	32
Total cores	48	96	192	384	768	1536
Speedup [Sequential time (s) ÷ Parallel Time (s)]	55/55 = 1.00	55/27 = 2.04	55/18 = 3.06	55/15 = 3.67	55/14 = 3.93	55/14 = 3.93

**Sequential Time :**  
**55 seconds**

# Efficiency

Efficiency vs Total Cores



The graph shows that as more cores are added, the system becomes less efficient. Each additional cores contributes less to overall performance due to increased coordination and communication overhead.

While more cores can speed up the system, they do not necessarily improve efficiency, as more resources are spent on managing the extra cores rather than productive work.

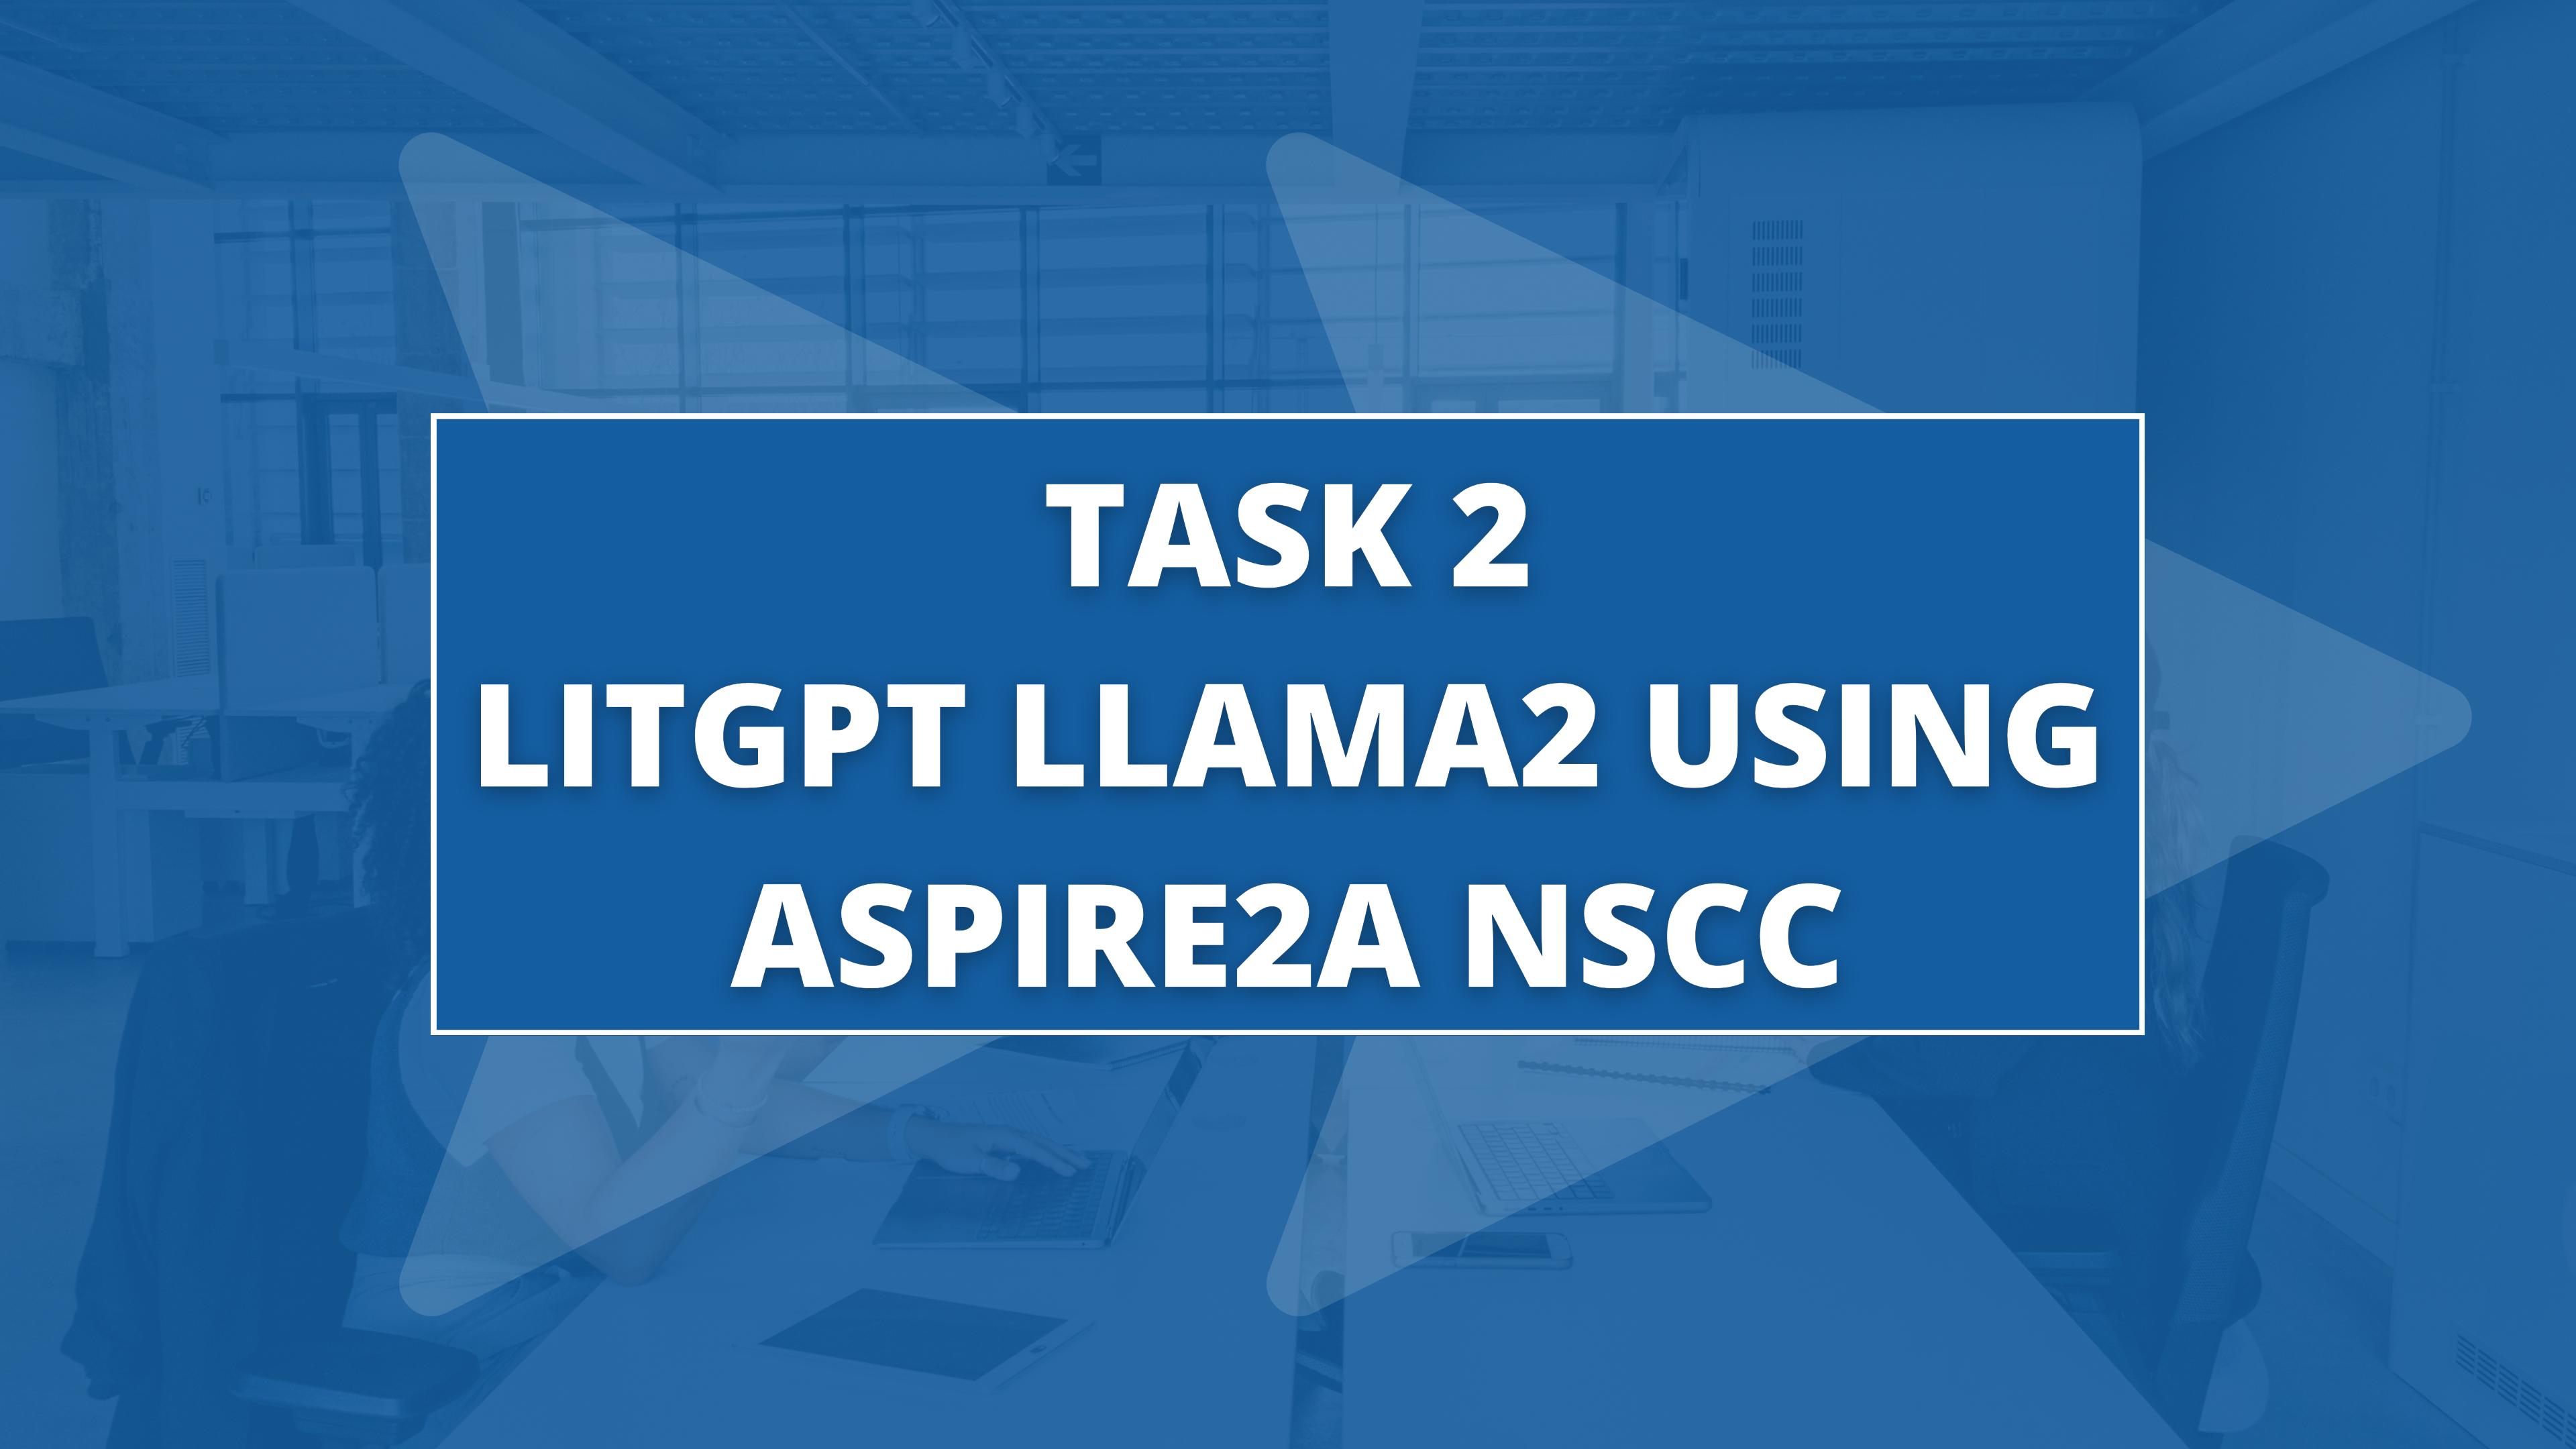
Number of nodes	1	2	4	8	16	32
Total cores	48	96	192	384	768	1536
Efficiency [Speedup/Total cores]	$1.00/48 = 0.02$	$2.037/96 = 0.021$	$3.056/192 = 0.016$	$3.667/384 = 0.010$	$3.929/768 = 0.005$	$3.929/1536 = 0.003$

# Challenges

- We aim to run simulations across multiple nodes, up to 32 in total. By scaling, the system resources like CPU, memory, and communication between nodes becomes difficult to handle.
- Initially, the performance increases with increasing cores. However, beyond a threshold, efficiency is reduced due to the dominating communication overhead between the different nodes.

# Conclusion

The results obtained from simulations demonstrate that an increased number of CPUs can significantly improve system performance by reducing execution time and enhancing simulation speed, especially when tasks are divided and processed in parallel. However, there are diminishing returns beyond a certain point, as coordination and communication overhead between nodes start to limit further improvements. This highlights the importance of an optimized balance between hardware scaling and system efficiency for maximum performance in high-performance computing environments.



# **TASK 2**

## **LITGPT LLAMA2 USING ASPIRE2A NSCC**

# Performance Metrics

01

**Training time (seconds)** - This measures how long it takes to complete the training.

# PBS Script Directives and Modules Loaded

```
#!/bin/bash
#PBS -P 50000032
#PBS -l walltime=00:01:00
#PBS -j oe
#PBS -M 214928@student.upm.edu.my,216638@student.upm.edu.my
#PBS -m abe

date
module purge
module load openmpi/4.1.2-hpe
module load libfabric/1.11.0.4.125
```

llama.sh

# PBS Value Initialization

<b>Number of nodes</b>	<b>Number of GPUs</b>	<b>Number of CPUs</b>	<b>Number of Epochs</b>	<b>Global Batch Size</b>	<b>Micro Batch Size</b>	<b>Max steps</b>
2	8	128	1	128	32	20

# PBS Commands Used

- qsub tuningllama.pbs - To submit job script (tuningllama.pbs)
- qstat - To display queue information
- qstat -q - To display current queue information

```
(base) [aa5323@gadi-login-03 run]$ qsub tuningllama.pbs  
126609921.gadi-pbs
```

```
(base) [aa5323@gadi-login-03 run]$ qstat  
Job id                               Name           User      Time Use S Queue  
-----  
126609921.gadi-pbs     tuningllama.pbs aa5323          0 Q gpuvolta-exec
```

```
(base) [aa5323@gadi-login-03 run]$ qstat -q  
  
server: gadi-pb
```

Queue	Memory	CPU	Time	Walltime	Node	Run	Que	Lm	State
normal	--	--	--	--	--	0	56	--	E R
normal-exec	--	--	--	--	1232	357	--	--	E R
express	--	--	--	--	--	0	0	--	E R
express-exec	--	--	--	--	--	2	1	--	E R
copyq	--	--	--	--	--	0	43	--	E R
copyq-exec	--	--	--	--	--	12	1	--	E R
gpuvolta	--	--	--	--	--	0	361	--	E R
gpuvolta-exec	--	--	--	--	--	339	424	--	E R
hugemem-exec	--	--	--	--	--	85	164	--	E R
hugemem	--	--	--	--	--	0	44	--	E R

# Execution Script

```
env
cat $PBS_NODEFILE

export NCCL_DEBUG=INFO
export NCCL_IB_DISABLE=1
export NCCL_NET_GDR_LEVEL=0
export NCCL_SHM_DISABLE=1

nvidia-smi

cmd="mpirun \
-wdir ${HOME}/scratch/workdir/llama \
-output-filename ${HOME}/run/output/${PBS_JOBNAME}.${PBS_JOBID} \
-map-by ppr:4:node -oversubscribe \
-report-bindings \
-x mpirun -mca btl ^ucx \
-mca coll_hcoll_enable 1 -mca coll_basic_priority 10 \
${HOME}/scratch/workdir/llama/litgpt.py312/bin/litgpt \
finetune_full \
${HOME}/scratch/workdir/llama/model/litgpt/meta-llama/Llama-2-7b-hf \
--data JSON --data.json_path ${HOME}/scratch/workdir/llama/dataset/alpaca1024 \
--out_dir ${HOME}/scratch/workdir/llama/out/finetune/full \
--config ${HOME}/scratch/workdir/llama/full.yaml \
--eval.final_validation=false \
--train.epochs=1 \
--devices=4 --num_nodes=2 \
--train.max_steps=${max_steps} \
```

tuningllama.pbs

- **NVIDIA-SMI**

Track gpu's utilization and performance

- **-x mpirun -mca btl ^ucx**

OpenMPI for communication backend, using Byte Transfer Layer (BTL) and exclude UCX

- **-mca coll\_hcoll\_enable 1**  
**-mca coll\_basic\_priority 10**

Set up collective algorithms such as enabling HCOLL and prioritize basic collective operations

# Submitting Job Commands

## Example:

```
[apacsc32@asp2a-login-nscc02 run]$ nodes=2 walltime=7201 \
> global_batch_size=128 micro_batch_size=32 max_steps=20 \
> bash -c \
> 'qsub -V \
> -l walltime=${walltime},select=${nodes}:ngpus=4 \
> -N llama.nodes${nodes}.GBS${global_batch_size}.MBS${micro_batch_size} \
> llama.sh'
8616103.pbs101
```

# Example of Output

NVIDIA-SMI 535.154.05			Driver Version: 535.154.05		CUDA Version: 12.2				
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC		
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.		
<hr/>									
0	NVIDIA A100-SXM4-40GB	On	00000000:03:00.0 Off	0MiB / 40960MiB	0%	Default	0		
N/A	43C	P0	52W / 400W	0MiB / 40960MiB	0%	Default	Disabled		
<hr/>									
1	NVIDIA A100-SXM4-40GB	On	00000000:41:00.0 Off	0MiB / 40960MiB	0%	Default	0		
N/A	42C	P0	51W / 400W	0MiB / 40960MiB	0%	Default	Disabled		
<hr/>									
2	NVIDIA A100-SXM4-40GB	On	00000000:81:00.0 Off	0MiB / 40960MiB	0%	Default	0		
N/A	43C	P0	51W / 400W	0MiB / 40960MiB	0%	Default	Disabled		
<hr/>									
3	NVIDIA A100-SXM4-40GB	On	00000000:C1:00.0 Off	0MiB / 40960MiB	0%	Default	0		
N/A	43C	P0	54W / 400W	0MiB / 40960MiB	0%	Default	Disabled		
<hr/>									
<b>Processes:</b>									
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage			
ID	ID								
<hr/>									
No running processes found									

llama.sh --> nvidia-smi output

# Example of Output

```
The longest sequence length in the train data is 512, the model's maximum sequence length is 512 and context length is 4096
Verifying settings ...
Epoch 1 | iter 1 step 1 | loss train: 1.417, val: n/a | iter time: 5379.64 ms (step)
Epoch 1 | iter 1 step 1 | loss train: 1.428, val: n/a | iter time: 5390.40 ms (step)
Epoch 1 | iter 2 step 2 | loss train: 1.627, val: n/a | iter time: 5081.73 ms (step)
Epoch 1 | iter 2 step 2 | loss train: 1.435, val: n/a | iter time: 5086.31 ms (step)
Epoch 1 | iter 3 step 3 | loss train: 1.331, val: n/a | iter time: 5092.07 ms (step)
Epoch 1 | iter 3 step 3 | loss train: 1.288, val: n/a | iter time: 5086.58 ms (step)
Epoch 1 | iter 4 step 4 | loss train: 1.010, val: n/a | iter time: 5083.41 ms (step)
Epoch 1 | iter 4 step 4 | loss train: 0.962, val: n/a | iter time: 5085.00 ms (step)
Epoch 2 | iter 5 step 5 | loss train: 0.880, val: n/a | iter time: 5258.94 ms (step)
Epoch 2 | iter 5 step 5 | loss train: 0.834, val: n/a | iter time: 5258.90 ms (step)

Training time: 28.07s
Memory used: 17.73 GB
Training time: 28.10s
Memory used: 17.73 GB
x1000c3s4b0n0:2898258:2898878 [3] NCCL INFO [Service thread] Connection closed by localRank 3
x1000c3s4b0n0:2898256:2898876 [1] NCCL INFO [Service thread] Connection closed by localRank 1
x1000c3s3b0n0:828731:829309 [3] NCCL INFO [Service thread] Connection closed by localRank 3
x1000c3s4b0n0:2898257:2898875 [2] NCCL INFO [Service thread] Connection closed by localRank 2
x1000c3s4b0n0:2898255:2898877 [0] NCCL INFO [Service thread] Connection closed by localRank 0
x1000c3s3b0n0:828729:829308 [1] NCCL INFO [Service thread] Connection closed by localRank 1
x1000c3s3b0n0:828730:829307 [2] NCCL INFO [Service thread] Connection closed by localRank 2
x1000c3s3b0n0:828728:829311 [0] NCCL INFO [Service thread] Connection closed by localRank 0
x1000c3s3b0n0:828731:832787 [3] NCCL INFO comm 0x14117290 rank 3 nranks 8 cudaDev 3 busId 81000 - Abort COMPLETE
x1000c3s4b0n0:2898258:2902251 [3] NCCL INFO comm 0x14a1e7b0 rank 7 nranks 8 cudaDev 3 busId 81000 - Abort COMPLETE
x1000c3s3b0n0:828730:832789 [2] NCCL INFO comm 0x1558c8c0 rank 2 nranks 8 cudaDev 2 busId 3000 - Abort COMPLETE
x1000c3s3b0n0:828729:832788 [1] NCCL INFO comm 0x1413d660 rank 1 nranks 8 cudaDev 1 busId c1000 - Abort COMPLETE
x1000c3s4b0n0:2898256:2902252 [1] NCCL INFO comm 0x152e6ff0 rank 5 nranks 8 cudaDev 1 busId c1000 - Abort COMPLETE
x1000c3s4b0n0:2898257:2902254 [2] NCCL INFO comm 0x143701b0 rank 6 nranks 8 cudaDev 2 busId 3000 - Abort COMPLETE
x1000c3s4b0n0:2898255:2902253 [0] NCCL INFO comm 0x14f65fa0 rank 4 nranks 8 cudaDev 0 busId 41000 - Abort COMPLETE
x1000c3s3b0n0:828728:832790 [0] NCCL INFO comm 0x161490e0 rank 0 nranks 8 cudaDev 0 busId 41000 - Abort COMPLETE
```

# LLAMA2 Results

## Baseline

<b>Number of nodes</b>	<b>Number of CPUs</b>	<b>Number of GPUs</b>	<b>Memory Used</b>	<b>Average Training Time (s)</b>
2	128	8	17.73GB	41.57

# LLAMA2 Results

## Improved script

```
cmd="mpirun \
-wdir ${HOME}/scratch/workdir/llama \
-output-filename ${HOME}/run/output/${PBS_JOBNAME}.${PBS_JOBID} \
-map-by ppr:4:node -oversubscribe \
-report-bindings \
-x mpirun -mca btl ^ucl \
-mca coll_hcoll_enable 1 -mca coll_basic_priority 10 \
```

Number of nodes	Number of CPUs	Number of GPUs	Memory Used	Average Training Time (s)
2	128	8	17.73GB	28.09

# LLAMA2 Results

## Improved script

```
export NCCL_DEBUG=INFO
export NCCL_IB_DISABLE=1
export NCCL_NET_GDR_LEVEL=0
export NCCL_SHM_DISABLE=1
```

- Enables detailed NCCL logging
- Disables InfiniBand, uses TCP/IP
- Disables GPU Direct RDMA, uses CPU/system memory.

Number of nodes	Number of CPUs	Number of GPUs	Memory Used	Average Training Time (s)
2	128	8	17.73GB	29.04

# LLAMA2 Results

## Fine tuning using LoRA

LoRA (Low-Rank Adaptation) is a technique that adapts pre-trained neural networks by introducing a small set of trainable low-rank matrices into the network. This approach reduces the number of parameters that need to be fine-tuned, making the adaptation process more efficient in terms of both computation and memory.

### **Observation:**

- Reduced training time
- Lower memory usage
- Maintained performance
- Flexibility and scalability

# Challenges

- In Gadi, we ran into some issues when trying to execute the job, which stopped us from getting any results from Gadi. It might be due to our building environment process. The disks ran out of memory while we were trying to install modules, thus we were not able to run the job successfully.
- In Aspire2A, achieving efficient resource utilization requires precise parameter adjustments (like batch size and learning rate) and an great understanding of MPI communication, which we are learning to navigate as we work on this project. As for the fine tuning using LoRA, we were not able to set it up on time due to our limited knowledge.

# Conclusion

This task that demanded in-depth understanding of HPC environments, parallel processing, and distributed machine learning that involves multi-faceted approach such as algorithmic advancements. Fortunately, this project successfully implemented optimizations that slightly reduced training time, such as adjustments to the communication backend.



# THANK YOU!

**UPM Team 3**

Universiti Putra Malaysia

